

Reasoning about Welfare-Affecting Capabilities in Concurrent Games

Yinfeng Li Emiliano Lorini

IRIT, CNRS, University of Toulouse

KR 2026

Strategic ability is not only about achieving goals

In multi-agent systems, a coalition's choices may also affect the welfare of other agents.

Classical ATL/CL asks

- Can coalition C ensure φ ?
- Can C keep φ true forever?

We also ask

- Can C benefit coalition B ?
- Can C harm coalition H ?
- Can it do so while ensuring φ ?

Who is a potential benefactor, and who is a danger?

Two Conceptual Ingredients

Harm and benefit

On the conceptual side, our notion of harm and benefit is closely related to the causal analysis in [Beckers et al., 2024]:

- harm: bringing about a bad outcome;
- benefit: bringing about a good outcome.

We adopt a consequentialist, non-comparative reading, where good and bad outcomes are represented through agents' preferences.

Danger

We follow a modal, non-probabilistic intuition for danger [Fassio, 2025, Garland, 2003, Peschard et al., 2022]:

- a coalition is dangerous if it has the potential to harm;
- potential means existence of a suitable strategy.

strategic capability [Alur et al., 2002, Goranko and van Drimmelen, 2006]

+

preference over outcomes [Li et al., 2025]

⇓

welfare-affecting capability

- 1 A semantics based on concurrent games enriched with agents' preferences over computations.
- 2 New ATL and CL modalities for welfare-affecting strategic capabilities:

$$\langle\langle C:B,H \rangle\rangle\gamma$$

“ C has a strategy that benefits B , harms H , and guarantees γ .”

- 3 Technical results:
 - sound and complete axiomatization for the until-free fragment;
 - polynomial embedding into standard ATL/CL with special atoms marking benefited and harmed agents;
 - tight satisfiability complexity: Exptime-complete for WA–ATL, Pspace-complete for WA–CL.

Models: Concurrent Games with Preferences

A model is a concurrent game model with regular preferences: We build on the recent framework of concurrent games with preferences.

$$M = (P, \mathcal{V})$$

where P contains:

- states and action profiles;
- a transition relation induced by action profiles;
- for each agent i and state s , a preference order


$$\preceq_{i,s}$$

over computations starting at s .

Regular preferences

We assume stability over time, non-flatness, and existence of best and worst computations. These assumptions support the strong notions of benefit and harm.

Running Example: Robots and Charging Rooms



- Robot 0 starts in the middle room.
- It wants to reach a charging station on the left or right.
- Robots 1 and 2 guard the right and left passages.
- Each guard can block or free its passage.

Preference profile

Robot 0 prefers computations where it eventually reaches a charging station.

Robots 1 and 2 prefer computations where robot 0 never reaches one.

Strong Benefit and Strong Harm

Let $\mathcal{O}_F(s, \mathcal{F}_C)$ be the computations compatible with strategy $\mathcal{F}_C$ from s .

Strongly beneficial for B

$$\mathcal{F}_C \in \text{StrSB}_M^C(s, B) \quad \text{iff} \quad \mathcal{O}_F(s, \mathcal{F}_C) \subseteq \text{Best}_{i,s} \text{ for every } i \in B.$$

Strongly harmful for H

$$\mathcal{F}_C \in \text{StrSH}_M^C(s, H) \quad \text{iff} \quad \mathcal{O}_F(s, \mathcal{F}_C) \subseteq \text{Worst}_{i,s} \text{ for every } i \in H.$$

Strong benefit/harm means guaranteeing best/worst outcomes for the affected coalition.

Fact

The coalition of robots $\{1, 2\}$ has a strategy that is harmful to robot 0 and beneficial to the guarding coalition.

$$\begin{aligned}\varphi, \psi &::= p \mid \neg\varphi \mid \varphi \wedge \psi \mid \langle\langle C:B,H \rangle\rangle\gamma \\ \gamma &::= X\varphi \mid G\varphi \mid \varphi U\psi\end{aligned}$$

Key modality

$$\langle\langle C:B,H \rangle\rangle\gamma$$

Coalition C has a strategy which is beneficial for B , harmful to H , and guarantees temporal property γ .

Standard ATL is recovered as the special case:

$$\langle\langle C \rangle\rangle\gamma \equiv \langle\langle C:\emptyset,\emptyset \rangle\rangle\gamma.$$

$$(M, s) \models \langle\langle C:B, H \rangle\rangle \gamma$$

iff there exists a strategy

$$\mathcal{F}_C \in \text{StrSB}_M^C(s, B) \cap \text{StrSH}_M^C(s, H)$$

such that every generated computation satisfies γ :

$$\forall \lambda \in \mathcal{O}_F(s, \mathcal{F}_C), \quad (M, \lambda) \models \gamma.$$

Three roles in one operator

- C : the acting coalition;
- B : the coalition benefited by the strategy;
- H : the coalition harmed by the strategy.

Potential Benefactor and Danger

The language gives compact definitions of the target notions.

Potential benefactor

$$\text{Benef}(C, B) \equiv \langle\langle C:B, \emptyset \rangle\rangle X\top.$$

Coalition C has a strategy beneficial for B .

Danger

$$\text{Danger}(C, H) \equiv \langle\langle C:\emptyset, H \rangle\rangle X\top.$$

Coalition C has a strategy harmful to H .

In the robot model:

$$(M_{rs}, s_0) \models \text{Benef}(\{1, 2\}, \{1, 2\}) \wedge \text{Danger}(\{1, 2\}, \{0\}).$$

We axiomatize the until-free, or safety, fragment:

$$\varphi ::= p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \langle\!\langle C:B,H \rangle\!\rangle X\varphi \mid \langle\!\langle C:B,H \rangle\!\rangle G\varphi.$$

Inference rules

$$\frac{\varphi, \quad \varphi \rightarrow \psi}{\psi} \quad (\text{MP}) \qquad \frac{\varphi \leftrightarrow \psi}{\langle\!\langle C:B,H \rangle\!\rangle X\varphi \leftrightarrow \langle\!\langle C:B,H \rangle\!\rangle X\psi} \quad (\text{RE}) \qquad \frac{\varphi}{\langle\!\langle C \rangle\!\rangle G\varphi} \quad (\text{Gen})$$

Axiom schemata

All tautologies of propositional logic	(TAUT)
$\neg \langle\langle C:B, H \rangle\rangle X \perp$	(NAAA)
$\langle\langle C:B, H \rangle\rangle X(\varphi \wedge \psi) \rightarrow \langle\langle C:B, H \rangle\rangle X\varphi$	(MG)
$\langle\langle C:B, H \rangle\rangle X\varphi \rightarrow \langle\langle C':B', H' \rangle\rangle X\varphi \quad \text{for } C \subseteq C', B \supseteq B', H \supseteq H'$	(MC)
$\langle\langle \emptyset \rangle\rangle X \top$	(NEI)
$(\langle\langle C:B, H \rangle\rangle X\varphi \wedge \langle\langle C':B', H' \rangle\rangle X\psi) \rightarrow \langle\langle C \cup C':B \cup B', H \cup H' \rangle\rangle X(\varphi \wedge \psi) \quad \text{for } C \cap C' = \emptyset$	(Sup)
$\langle\langle C:B, H \rangle\rangle X(\varphi \vee \psi) \rightarrow (\langle\langle C:B, H \rangle\rangle X\varphi \vee \langle\langle \text{AGT}:B, H \rangle\rangle X\psi)$	(Cro)
$\langle\langle C:B, H \rangle\rangle G\varphi \leftrightarrow \langle\langle C:B, H \rangle\rangle X(\varphi \wedge \langle\langle C:B, H \rangle\rangle G\varphi)$	(FP _G)
$\langle\langle \emptyset \rangle\rangle G(\psi \rightarrow \langle\langle C:B, H \rangle\rangle X(\varphi \wedge \psi)) \rightarrow \langle\langle \emptyset \rangle\rangle G(\psi \rightarrow \langle\langle C:B, H \rangle\rangle G\varphi)$	(GFP _G)
$\langle\langle \text{AGT}:\emptyset, \{a\} \rangle\rangle X \top$	(WFP)
$\langle\langle \text{AGT}:\{a\}, \emptyset \rangle\rangle X \top$	(CWFP)
$\neg \langle\langle C:B, H \rangle\rangle X \top \quad \text{for } B \cap H \neq \emptyset$	(BHC)

Theorem

For every $\varphi \in \mathcal{L}_{\text{WA-ATL-u}}$:

φ is valid over CGMRPs iff $\vdash \varphi$.

- The proof uses canonical models built under a formula context.
- The states in the canonical models consist of (*finite*) maximally consistent sets of formulas in the formula context, and the labels of benefited and harmed groups.

Theorem

For every $\varphi \in \mathcal{L}_{\text{WA-ATL}}$:

φ is satisfiable iff φ is satisfiable on a pointed tree-like CGMRP.

- Tree-like models have a unique root, unique predecessors, and no cycles.
- This lets us localize welfare labels along generated computations.
- Stability of preferences keeps those labels coherent over time.

Embedding into Standard ATL

Extend the atomic vocabulary with:

$$\heartsuit_i \quad \text{and} \quad \spadesuit_i$$

read as “agent i is benefited” and “agent i is harmed”.

Translation:

$$tr(\langle\langle C:B,H \rangle\rangle X\varphi) = \langle\langle C \rangle\rangle X(\heartsuit B \wedge \spadesuit H \wedge tr(\varphi)),$$

where

$$\heartsuit B = \bigwedge_{i \in B} \heartsuit_i, \quad \spadesuit H = \bigwedge_{i \in H} \spadesuit_i.$$

Similar clauses handle G and U.

Embedding Theorem and Complexity

Polynomial embedding

Satisfiability for WA–ATL reduces polynomially to satisfiability for standard ATL, with additional constraints ensuring:

best labels exist, worst labels exist, and no state is both for the same agent.

Complexity

Language	Target	Satisfiability
WA–ATL	ATL	Exptime-complete
WA–CL	CL	Pspace-complete

The bounds are tight because ATL and CL are special cases.

Summary and Outlook

- Strategic agency can be welfare-affecting: agents can benefit or endanger others through their strategies.
- WA–ATL internalizes this idea in the object language:

$$\langle\langle C:B, H \rangle\rangle\gamma.$$

- Potential benefactor and danger become definable notions.
- The logic remains technically well-behaved:
 - complete axiomatization for the safety fragment;
 - polynomial embedding into standard ATL/CL with welfare-label atoms;
 - tight complexity bounds.

Next steps

Full axiomatization with until; probabilistic risk-based danger; rational harmful strategies and strong danger.

References I

- Rajeev Alur, Thomas A Henzinger, and Orna Kupferman. Alternating-time temporal logic. *Journal of the ACM*, 49(5):672–713, 2002.
- Sander Beckers, Hana Chockler, and Joseph Y. Halpern. A causal analysis of harm. *Minds and Machines*, 34(3):1–24, 2024.
- Davide Fassio. On danger. *Philosophy and Phenomenological Research*, 2025. doi: 10.1111/phpr.70036. Version of Record online: 08 August 2025.
- David Garland. The rise of risk. In Richard V. Ericson and Aaron Doyle, editors, *Risk and Morality*, pages 48–86. University of Toronto Press, 2003.
- Valentin Goranko and Govert van Drimmelen. Complete axiomatization and decidability of alternating-time temporal logic. *Theoretical Computer Science*, 353:93–117, 2006.
- Yinfeng Li, Emiliano Lorini, and Munyque Mittelmann. Rational capability in concurrent games. In *Proceedings of the 24th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2025)*, pages 1309–1317. International Foundation for Autonomous Agents and Multiagent Systems / ACM, 2025.
- Isabelle Peschard, Yann Benétreau-Dupin, and Christopher Wessels. *Philosophy and Science of Risk: An Introduction*. Routledge, New York, 2022.

Thank you

Questions?